

Problem Statement

Polo Ralph Lauren receives sales and return data from multiple channels — **POS**, **E-commerce**, and **Third-Party partners** — in separate raw CSV files.

These files:

- Arrive in **different currencies** and are stored in Amazon S3
- Have to be **reconciled** to produce a unified view of sales
- Need to be **kept updated automatically** without manual loading

The business requires:

- A **single fact table** combining all channels
 - **Amounts standardized to USD** for consistent reporting
 - **Refund transactions identified** for tracking returns
 - **Automated daily updates** when new files arrive in S3
-

Solution Implemented (AWS + Snowflake ELT)

1 Data Storage & Ingestion

- Created an **S3 bucket** `omni-recon-data` with folders:

```
raw/pos/  
raw/ecommerce/  
raw/thirdparty/  
raw/reference/
```

- Uploaded channel-specific CSV files + FX rate file to their respective folders.
 - Created a **Snowflake storage integration** with AWS IAM role & external stage `@omni_stage` pointing to the S3 bucket.
-

2 Automatic Loading (Snowpipe)

- Created **4 Snowpipes**:
 - `pos_pipe` → loads from `raw/pos/` → `stage_pos_sales`
 - `ecommerce_pipe` → loads from `raw/ecommerce/` → `stage_ecommerce_sales`
 - `thirdparty_pipe` → loads from `raw/thirdparty/` → `stage_thirdparty_sales`
 - `fx_rates_pipe` → loads from `raw/reference/` → `stage_fx_rates`
 - Configured them to **auto-ingest** new files from S3 into staging tables.
-

3 Staging Table Structures

All sales staging tables (POS , E-commerce , Third-Party):

```

transaction_id STRING
channel STRING
store_id STRING
product_id STRING
currency STRING
amount_local FLOAT
refund_flag BOOLEAN
quantity INT
datetime_local TIMESTAMP

```

FX rates staging table:

```

date DATE
currency STRING
exchange_rate FLOAT

```

4 Transformations Implemented

The transformation logic inside Snowflake creates the `sales_reconciliation_fact` table by:

1. Merging all sales channels

```

SELECT * FROM stage_pos_sales UNION ALL SELECT * FROM stage_ecommerce_sales
UNION ALL SELECT * FROM stage_thirdparty_sales

```

2. Joining with FX rates to standardize amounts in USD

```

ROUND(s.amount_local / f.exchange_rate, 2) AS amount_usd

```

3. Retaining refund_flag to identify returns

```

s.refund_flag

```

4. Extracting date dimensions for reporting

```

YEAR(s.datetime_local) AS year, MONTH(s.datetime_local) AS month

```

5. Deduplicating on load (incremental)

```

WHERE NOT EXISTS ( SELECT 1 FROM sales_reconciliation_fact fact WHERE
fact.transaction_id = s.transaction_id AND fact.channel = s.channel )

```

5 Automation (Snowflake Task)

- Created a **daily scheduled Snowflake Task** `refresh_sales_fact` :
 - Runs at **2 AM UTC**.
 - Inserts only **new transactions** from staging tables into `sales_reconciliation_fact`.
 - Leaves existing rows untouched → **incremental loads**, not full rebuilds.
-

6 Current Workflow

1. **New file lands in S3** → Snowpipe detects and loads into staging table.
2. **Daily at 2 AM UTC** → Task runs transformation SQL → Inserts only new rows into fact table.
3. `sales_reconciliation_fact` always contains **all channels, all currencies converted to USD, refunds tracked, with Year/Month fields** for easy reporting.